

cm006: dplyr Exercise

Optional, but recommended startup:

1. Change the file output to both html and md *documents* (not notebook).
2. `knit` the document.
3. Stage and commit the rmd, and knitted documents.

Intro to dplyr syntax

Load the `gapminder` and `tidyverse` packages. Hint: `suppressPackageStartupMessages()!` - This loads `dplyr`, too.

```
# load your packages here:
library(gapminder)
library(tidyverse)
```

`select()` (8 min)

1. Make a data frame containing the columns `year`, `lifeExp`, `country` from the `gapminder` data, in that order.

```
select(gapminder, year, lifeExp, country)
```

```
## # A tibble: 1,704 x 3
##   year lifeExp country
##   <int>   <dbl> <fct>
## 1  1952    28.8 Afghanistan
## 2  1957    30.3 Afghanistan
## 3  1962    32.0 Afghanistan
## 4  1967    34.0 Afghanistan
## 5  1972    36.1 Afghanistan
## 6  1977    38.4 Afghanistan
## 7  1982    39.9 Afghanistan
## 8  1987    40.8 Afghanistan
## 9  1992    41.7 Afghanistan
## 10 1997    41.8 Afghanistan
## # ... with 1,694 more rows
```

```
#Choose variables from a tbl. select() keeps only the variables you mention#
```

2. Select all variables, from `country` to `lifeExp`.

```
# This will work:
select(gapminder, country, continent, year, lifeExp)
```

```
## # A tibble: 1,704 x 4
##   country      continent  year lifeExp
##   <fct>       <fct>    <int>  <dbl>
## 1 Afghanistan Asia      1952   28.8
## 2 Afghanistan Asia      1957   30.3
## 3 Afghanistan Asia      1962   32.0
## 4 Afghanistan Asia      1967   34.0
## 5 Afghanistan Asia      1972   36.1
## 6 Afghanistan Asia      1977   38.4
## 7 Afghanistan Asia      1982   39.9
## 8 Afghanistan Asia      1987   40.8
## 9 Afghanistan Asia      1992   41.7
## 10 Afghanistan Asia      1997   41.8
## # ... with 1,694 more rows
```

```
# Better way:
select(gapminder, country:lifeExp) #same function as [from:to] in vector
```

```
## # A tibble: 1,704 x 4
##   country      continent  year lifeExp
##   <fct>       <fct>    <int>  <dbl>
## 1 Afghanistan Asia      1952   28.8
## 2 Afghanistan Asia      1957   30.3
## 3 Afghanistan Asia      1962   32.0
## 4 Afghanistan Asia      1967   34.0
## 5 Afghanistan Asia      1972   36.1
## 6 Afghanistan Asia      1977   38.4
## 7 Afghanistan Asia      1982   39.9
## 8 Afghanistan Asia      1987   40.8
## 9 Afghanistan Asia      1992   41.7
## 10 Afghanistan Asia      1997   41.8
## # ... with 1,694 more rows
```

3. Select all variables, except lifeExp.

```
select(gapminder, lifeExp)
```

```
## # A tibble: 1,704 x 1
##   lifeExp
##   <dbl>
## 1   28.8
## 2   30.3
## 3   32.0
## 4   34.0
## 5   36.1
## 6   38.4
## 7   39.9
## 8   40.8
## 9   41.7
## 10  41.8
## # ... with 1,694 more rows
```

4. Put continent first. Hint: use the `everything()` function.

```
select(gapminder, continent, 'everything'())
```

```
## # A tibble: 1,704 x 6
##   continent country      year lifeExp      pop gdpPercap
##   <fct>      <fct>      <int>  <dbl>    <int>    <dbl>
## 1 Asia      Afghanistan 1952    28.8  8425333    779.
## 2 Asia      Afghanistan 1957    30.3  9240934    821.
## 3 Asia      Afghanistan 1962    32.0 10267083    853.
## 4 Asia      Afghanistan 1967    34.0 11537966    836.
## 5 Asia      Afghanistan 1972    36.1 13079460    740.
## 6 Asia      Afghanistan 1977    38.4 14880372    786.
## 7 Asia      Afghanistan 1982    39.9 12881816    978.
## 8 Asia      Afghanistan 1987    40.8 13867957    852.
## 9 Asia      Afghanistan 1992    41.7 16317921    649.
## 10 Asia     Afghanistan 1997    41.8 22227415    635.
## # ... with 1,694 more rows
```

5. Rename continent to cont.

```
# compare
select(gapminder, cont = continent, everything()) # select contitnent and keep everthing else
```

```
## # A tibble: 1,704 x 6
##   cont country      year lifeExp      pop gdpPercap
##   <fct> <fct>      <int>  <dbl>    <int>    <dbl>
## 1 Asia  Afghanistan 1952    28.8  8425333    779.
## 2 Asia  Afghanistan 1957    30.3  9240934    821.
## 3 Asia  Afghanistan 1962    32.0 10267083    853.
## 4 Asia  Afghanistan 1967    34.0 11537966    836.
## 5 Asia  Afghanistan 1972    36.1 13079460    740.
## 6 Asia  Afghanistan 1977    38.4 14880372    786.
## 7 Asia  Afghanistan 1982    39.9 12881816    978.
## 8 Asia  Afghanistan 1987    40.8 13867957    852.
## 9 Asia  Afghanistan 1992    41.7 16317921    649.
## 10 Asia  Afghanistan 1997    41.8 22227415    635.
## # ... with 1,694 more rows
```

```
select(gapminder, country, gdpPercap:continent) # display `country` first, and reverse the order for re.
```

```
## # A tibble: 1,704 x 6
##   country      gdpPercap      pop lifeExp year continent
##   <fct>      <dbl>    <int>  <dbl> <int> <fct>
## 1 Afghanistan    779.  8425333    28.8  1952 Asia
## 2 Afghanistan    821.  9240934    30.3  1957 Asia
## 3 Afghanistan    853. 10267083    32.0  1962 Asia
## 4 Afghanistan    836. 11537966    34.0  1967 Asia
## 5 Afghanistan    740. 13079460    36.1  1972 Asia
## 6 Afghanistan    786. 14880372    38.4  1977 Asia
## 7 Afghanistan    978. 12881816    39.9  1982 Asia
## 8 Afghanistan    852. 13867957    40.8  1987 Asia
## 9 Afghanistan    649. 16317921    41.7  1992 Asia
## 10 Afghanistan    635. 22227415    41.8  1997 Asia
## # ... with 1,694 more rows
```

```
rename(gapminder, cont = continent) # rename the continent, keep the original order
```

```
## # A tibble: 1,704 x 6
##   country    cont  year lifeExp      pop gdpPercap
##   <fct>      <fct> <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia   1952   28.8  8425333    779.
## 2 Afghanistan Asia   1957   30.3  9240934    821.
## 3 Afghanistan Asia   1962   32.0 10267083    853.
## 4 Afghanistan Asia   1967   34.0 11537966    836.
## 5 Afghanistan Asia   1972   36.1 13079460    740.
## 6 Afghanistan Asia   1977   38.4 14880372    786.
## 7 Afghanistan Asia   1982   39.9 12881816    978.
## 8 Afghanistan Asia   1987   40.8 13867957    852.
## 9 Afghanistan Asia   1992   41.7 16317921    649.
## 10 Afghanistan Asia   1997   41.8 22227415    635.
## # ... with 1,694 more rows
```

arrange() (8 min)

1. Order by year.

```
arrange(gapminder, year)
```

```
## # A tibble: 1,704 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Albania     Europe    1952   55.2  1282697   1601.
## 3 Algeria     Africa    1952   43.1  9279525   2449.
## 4 Angola      Africa    1952   30.0  4232095   3521.
## 5 Argentina   Americas  1952   62.5 17876956   5911.
## 6 Australia   Oceania   1952   69.1  8691212  10040.
## 7 Austria     Europe    1952   66.8  6927772   6137.
## 8 Bahrain     Asia      1952   50.9   120447   9867.
## 9 Bangladesh  Asia      1952   37.5 46886859    684.
## 10 Belgium     Europe    1952    68   8730405   8343.
## # ... with 1,694 more rows
```

2. Order by year, in descending order.

```
arrange(gapminder, desc(year))
```

```
## # A tibble: 1,704 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      2007   43.8 31889923    975.
## 2 Albania     Europe    2007   76.4  3600523   5937.
## 3 Algeria     Africa    2007   72.3 33333216   6223.
## 4 Angola      Africa    2007   42.7 12420476   4797.
## 5 Argentina   Americas  2007   75.3 40301927  12779.
```

```
## 6 Australia Oceania 2007 81.2 20434176 34435.
## 7 Austria Europe 2007 79.8 8199783 36126.
## 8 Bahrain Asia 2007 75.6 708573 29796.
## 9 Bangladesh Asia 2007 64.1 150448339 1391.
## 10 Belgium Europe 2007 79.4 10392226 33693.
## # ... with 1,694 more rows
```

3. Order by year, then by life expectancy.

```
arrange(gapminder, year, lifeExp)
```

```
## # A tibble: 1,704 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>   <int>   <dbl>
## 1 Afghanistan Asia      1952    28.8 8425333    779.
## 2 Gambia      Africa    1952     30  284320     485.
## 3 Angola      Africa    1952    30.0 4232095   3521.
## 4 Sierra Leone Africa    1952    30.3 2143249     880.
## 5 Mozambique  Africa    1952    31.3 6446316     469.
## 6 Burkina Faso Africa    1952    32.0 4469979     543.
## 7 Guinea-Bissau Africa    1952    32.5  580653     300.
## 8 Yemen, Rep. Asia      1952    32.5 4963829     782.
## 9 Somalia     Africa    1952    33.0 2526994   1136.
## 10 Guinea     Africa    1952    33.6 2664249     510.
## # ... with 1,694 more rows
```

Piping, %>% (8 min)

Note: think of %>% as the word “then”!

Demonstration:

Here I want to combine `select()` Task 1 with `arrange()` Task 3.

This is how I could do it by *nesting* the two function calls:

```
# Nesting function calls can be hard to read
# the inner function will be excuted first, then is the outer function.
arrange(select(gapminder, year, lifeExp, country), year, lifeExp)
```

Now using with pipes:

```
# alter the below to include 2 "pipes"
gapminder %>% # select(.data, ...), gapminder here will be the data
  select(year, lifeExp, country) %>%
  arrange(year, lifeExp)
```

```
## # A tibble: 1,704 x 3
##   year lifeExp country
##   <int>   <dbl> <fct>
## 1 1952    28.8 Afghanistan
## 2 1952     30 Gambia
## 3 1952    30.0 Angola
```

```
## 4 1952 30.3 Sierra Leone
## 5 1952 31.3 Mozambique
## 6 1952 32.0 Burkina Faso
## 7 1952 32.5 Guinea-Bissau
## 8 1952 32.5 Yemen, Rep.
## 9 1952 33.0 Somalia
## 10 1952 33.6 Guinea
## # ... with 1,694 more rows
```

Resume lecture

Return to guide at section 6.7.

filter() (10 min)

1. Only take data with population greater than 100 million.

```
gapminder %>%
  filter(pop>10000000)
```

```
## # A tibble: 659 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1962   32.0 10267083    853.
## 2 Afghanistan Asia      1967   34.0 11537966    836.
## 3 Afghanistan Asia      1972   36.1 13079460    740.
## 4 Afghanistan Asia      1977   38.4 14880372    786.
## 5 Afghanistan Asia      1982   39.9 12881816    978.
## 6 Afghanistan Asia      1987   40.8 13867957    852.
## 7 Afghanistan Asia      1992   41.7 16317921    649.
## 8 Afghanistan Asia      1997   41.8 22227415    635.
## 9 Afghanistan Asia      2002   42.1 25268405    727.
## 10 Afghanistan Asia      2007   43.8 31889923    975.
## # ... with 649 more rows
```

2. Your turn: of those rows filtered from step 1., only take data from Asia.

```
gapminder %>%
  filter(pop>10000000, continent == "Asia")
```

```
## # A tibble: 232 x 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1962   32.0 10267083    853.
## 2 Afghanistan Asia      1967   34.0 11537966    836.
## 3 Afghanistan Asia      1972   36.1 13079460    740.
## 4 Afghanistan Asia      1977   38.4 14880372    786.
## 5 Afghanistan Asia      1982   39.9 12881816    978.
## 6 Afghanistan Asia      1987   40.8 13867957    852.
## 7 Afghanistan Asia      1992   41.7 16317921    649.
```

```
## 8 Afghanistan Asia      1997    41.8 22227415    635.
## 9 Afghanistan Asia      2002    42.1 25268405    727.
## 10 Afghanistan Asia     2007    43.8 31889923    975.
## # ... with 222 more rows
```

3. Repeat 2, but take data from countries Brazil, and China.

```
gapminder %>%
  filter(pop>10000000,
         country == 'Brazil'|country == 'China')
```

```
## # A tibble: 24 x 6
##   country continent  year lifeExp      pop gdpPercap
##   <fct>    <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Brazil  Americas   1952   50.9 56602560   2109.
## 2 Brazil  Americas   1957   53.3 65551171   2487.
## 3 Brazil  Americas   1962   55.7 76039390   3337.
## 4 Brazil  Americas   1967   57.6 88049823   3430.
## 5 Brazil  Americas   1972   59.5 100840058  4986.
## 6 Brazil  Americas   1977   61.5 114313951  6660.
## 7 Brazil  Americas   1982   63.3 128962939  7031.
## 8 Brazil  Americas   1987   65.2 142938076  7807.
## 9 Brazil  Americas   1992   67.1 155975974  6950.
## 10 Brazil Americas   1997   69.4 168546719  7958.
## # ... with 14 more rows
```

mutate() (10 min)

Let's get:

- GDP by multiplying GDP per capita with population, and
- GDP in billions, named (gdpBill), rounded to two decimals.

```
gapminder %>%
  mutate(gdpBill = round((gdpPercap*pop/1000000), digits=2))
```

```
## # A tibble: 1,704 x 7
##   country    continent  year lifeExp      pop gdpPercap gdpBill
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.    657.
## 2 Afghanistan Asia      1957   30.3  9240934    821.    759.
## 3 Afghanistan Asia      1962   32.0 10267083    853.    876.
## 4 Afghanistan Asia      1967   34.0 11537966    836.    965.
## 5 Afghanistan Asia      1972   36.1 13079460    740.    968.
## 6 Afghanistan Asia      1977   38.4 14880372    786.   1170.
## 7 Afghanistan Asia      1982   39.9 12881816    978.   1260.
## 8 Afghanistan Asia      1987   40.8 13867957    852.   1182.
## 9 Afghanistan Asia      1992   41.7 16317921    649.   1060.
## 10 Afghanistan Asia      1997   41.8 22227415    635.   1412.
## # ... with 1,694 more rows
```

Notice the backwards compatibility! No need for loops!

Try the same thing, but with `transmute` (drops all other variables).

```
gapminder %>%
  transmute(gdpBill = gdpPercap*pop/1000000)
```

```
## # A tibble: 1,704 x 1
##   gdpBill
##   <dbl>
## 1    657.
## 2    759.
## 3    876.
## 4    965.
## 5    968.
## 6   1170.
## 7   1260.
## 8   1182.
## 9   1060.
## 10  1412.
## # ... with 1,694 more rows
```

The `if_else` function is useful for changing certain elements in a data frame.

Example: Suppose Canada's 1952 life expectancy was mistakenly entered as 68.8 in the data frame, but is actually 70. Fix it using `if_else` and `mutate`.

```
gapminder %>%
  mutate(lifeExp = if_else(country == 'Canada' & year == 1952, 70, lifeExp))
```

```
## # A tibble: 1,704 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 1,694 more rows
```

Your turn: Make a new column called `cc` that pastes the country name followed by the continent, separated by a comma. (Hint: use the `paste` function with the `sep=","` argument).

```
gapminder %>%
  mutate(cc = paste(country, continent, sep = ", "))
```

```
## # A tibble: 1,704 x 7
##   country    continent  year lifeExp      pop gdpPercap cc
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl> <fct>
```



```
##      <fct>      <fct>      <int>      <dbl>      <int>      <dbl> <chr>
## 1 Afghanistan Asia      1952      28.8  8425333    779. Afghanistan, Asia
## 2 Afghanistan Asia      1957      30.3  9240934    821. Afghanistan, Asia
## 3 Afghanistan Asia      1962      32.0 10267083    853. Afghanistan, Asia
## 4 Afghanistan Asia      1967      34.0 11537966    836. Afghanistan, Asia
## 5 Afghanistan Asia      1972      36.1 13079460    740. Afghanistan, Asia
## 6 Afghanistan Asia      1977      38.4 14880372    786. Afghanistan, Asia
## 7 Afghanistan Asia      1982      39.9 12881816    978. Afghanistan, Asia
## 8 Afghanistan Asia      1987      40.8 13867957    852. Afghanistan, Asia
## 9 Afghanistan Asia      1992      41.7 16317921    649. Afghanistan, Asia
## 10 Afghanistan Asia      1997      41.8 22227415    635. Afghanistan, Asia
## # ... with 1,694 more rows
```

These functions we've seen are called **vectorized functions**.

git stuff (Optional)

Knit, commit, push!

Bonus Exercises

If there's time remaining, we'll practice with these three exercises. I'll give you 1 minute for each, then we'll go over the answer.

1. Take all countries in Europe that have a GDP per capita greater than 10000, and select all variables except `gdpPercap`. (Hint: use `-`).
2. Take the first three columns, and extract the names.
3. Of the `iris` data frame, take all columns that start with the word "Petal".
 - Hint: take a look at the "Select helpers" documentation by running the following code:
`?tidyselect::select_helpers`.
4. Convert the population to a number in billions.
5. Filter the rows of the `iris` dataset for `Sepal.Length >= 4.6` and `Petal.Width >= 0.5`.

Exercises 3. and 5. are from r-exercises.